

AI & Automation big bets

# **Where we are, how we are resourcing the work, and how we propose to sequence and fund it**

Discussion document for Enrique and Jaime | Friday, September 4, 2026

Also serves the ELT session on Tuesday, September 8 (pipeline tracking, POC progress, CEO/CFO follow-ups, sequencing approach)

Enterprise Transformation Office | Presented by Anshu

PRELIMINARY - WORKING DRAFT. Figures are pre-S2 submission and have not yet been reviewed with Finance or sponsors; numeric detail is held in speaker notes only.

# Today is a level-set ahead of S2 pencils-down: orient you on where the pipeline stands and what delivery will take, so nothing next week is a surprise

## 1 Level-set on where we are

- Snapshot of the big-bet pipeline we are tracking and its S2 status
- Where Wave 1 org changes are affecting baselines and sizing
- What has been engaged with technology, Finance and HR so far

## 2 Share resourcing and propose sequencing

- Current working model and the typical build-and-run team we are sizing against
- Technology resourcing capacity view (Srini's framework, to be inserted)
- Proposed criteria for sequencing and prioritization across levers

## 3 Preview investment themes, not dollars

- Shape and magnitude of the investment we expect to request
- Where the savings actually come from, and what that implies
- Accounting treatment and funding-source questions being worked with Finance

### What today is

- An orientation and a preview of the asks we expect to make after S2 closes
- A chance to align on method (sequencing, resourcing, funding) before numbers land

### What today is not


- A request to approve investment cases: S2 submissions close this week and are reviewed with Finance and Jaime next week
- An initiative-level readout: numbers will come to you consolidated after sponsor and Finance alignment

Next two weeks: S2 pencils-down (this week) > Finance review > Jaime pressure-test sessions (next week) > ELT update (Tuesday) > S3 in ~2.5 weeks


# Executive summary: the big-bet pipeline is converging in S2, proven use cases can move first, and delivery will need dedicated capacity and a clear funding path

|   |                           |                                                                                                                                                                                                                                                                                                                            |
|---|---------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1 | <b>Pipeline</b>           | We are tracking roughly 15-16 big-bet initiatives across five areas (servicing, seller productivity, SDLC, PDLC, risk and compliance). Most are in S2 sizing; the remainder of the automation lever is bottom-up initiatives already in Wave.                                                                              |
| 2 | <b>Where we are</b>       | S1 sizing was shared with sponsors in mid-August. Since then we have run the three seller-productivity initiatives to ground with the sales working team, aligned the servicing model with ops leaders, and brought Finance and HR into the business cases. S2 pencils-down is this week; S3 is roughly 2.5 weeks out.     |
| 3 | <b>Org change impacts</b> | Wave 1 actions moved the ground under the sizing: baselines (servicing headcount, product and engineering ratios) are being re-cut, and several SMEs and finance partners who supported S1 are no longer here. We are re-baselining rather than carrying stale assumptions.                                                |
| 4 | <b>Resourcing</b>         | Today the work is carried by a core ETO team with consulting support, workstream leads, tech POCs, Finance and HR. Delivery will need dedicated, funded build and run teams; we are sizing against a standard team archetype and phase plan and will bring Srin's tech capacity view.                                      |
| 5 | <b>Sequencing</b>         | We propose to sequence on confidence, timing of value, investment intensity, dependencies across levers, and execution capacity, and to tier investment to confidence: proven use cases (servicing, SDLC) move left now; less proven ones (for example PDLC) get design-and-pilot funding gated on evidence.               |
| 6 | <b>Investment</b>         | Investment will be material and front-loaded into build, with recurring run cost (inference, run teams) as the swing factor. Within servicing most savings come from contact avoidance; how seller-productivity value is taken is an open sponsor and Finance question. Sized numbers follow the Finance review next week. |

# We are mid-way through S2 with S3 roughly 2.5 weeks out; each gate raises confidence from theoretical sizing to piloted evidence, so estimates will keep moving until S4



## Who has been engaged in the sizing so far, and who is still to be engaged



Implication: numbers shown between now and S4 are directional by design. We re-baseline at each gate on the organization we actually have after Wave 1, rather than freezing early estimates.

# We are tracking roughly 15-16 big-bet initiatives across five areas; most are in S2 sizing, and the profile is weighted to proven use cases in servicing and SDLC


Status is the ETO read as of Sep 2, to be confirmed with sponsors before the ELT session; counts to be confirmed from the Wave export; initiative-level investment is intentionally not shown

| Big-bet area        | # (approx.)            | What it covers                                                                                                             | S2 status                                                                                           | Status                                                                                | ETO commentary                                                                                                                 |
|---------------------|------------------------|----------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------|
| Servicing           | 5-6                    | Contact avoidance; enhanced chat and self-service; agentic voice; skill-based routing; workforce management; non-FTE costs | Sizing aligned with ops leaders Sep 2; submission tracking to Friday                                |    | Best-understood, well-trodden use cases; savings dominated by contact avoidance. Top-of-funnel headroom flagged for S3         |
| Seller productivity | 3 fast-tracked         | Underwriting, pricing, SMB long tail (plus sales integration cost-out running as a bottom-up initiative)                   | Pressure-tested with sales working team; sponsor (Susan) review Sep 3; Finance review pending       |    | Cost-out versus reinvestment treatment is an open sponsor and Finance question; two components not yet sized pending SME input |
| SDLC                | 1 (confirm)            | AI-assisted software development lifecycle across engineering                                                              | In S2 sizing with tech POCs                                                                         |    | Well-trodden, high-value use case; recurring inference cost is the swing factor and efficiency work is under way               |
| PDLC                | 1 parent + BU children | Standardized product-development lifecycle with AI tooling for product, design and adjacent roles                          | Baseline being re-cut post Wave 1 and role-ratio actions; BU placeholders zeroed pending validation |    | Centrally managed to avoid double counting; value ramps late (design and pilots first)                                         |
| Risk and compliance | 3-4                    | Financial crimes, collections, Helix platform, AML (small)                                                                 | Collections submitted for S2; financial-crimes cases in Finance review; Helix case on its own track |  | Smaller cases are straightforward; the Helix case needs its execution plan reviewed before it is sequenced with the rest       |


 On track for S2       Open dependency or validation before S3       At risk (none today)

*Bottom-up automation initiatives outside the big bets continue through the standard S2 process in Wave and are not shown here.*

# Wave 1 org actions changed the ground under the sizing: baselines and the people who built them have moved, so we are re-baselining rather than carrying stale assumptions



## Where we see it today



Ask: acknowledge that S2 numbers reflect a moving organization, and that some re-validation time is the cost of not double counting.

# Today a core ETO team carries the work with technology, Finance and HR; delivery will need dedicated, funded build and run teams sized to a typical team model

Placeholder: Srini's technology resourcing framework (already shared with Jaime) to be inserted as the following page once obtained

## How the work is resourced today

### ETO core team

Program management, pipeline tracking in Wave, sponsor and ELT cadence; McKinsey support on sizing, models and business cases

### Workstream leads and sponsors

Own the business cases; provide SMEs, data and sign-off at each gate


### Technology points of contact

Roles per build, phase durations, parallelization, early view of which roles could be capitalized

### Finance, HR, Controllershship

Business-case review, headcount counting, accounting treatment

## What delivery will need: a typical build-and-run team per initiative



### Build team (one-time)

- Product management, engineering, QA, DevOps, UX at a blended cost per person
- Teams and duration set per phase with tech POCs; some phases run in parallel
- Which roles can be capitalized is being worked with Finance and Controllershship

### Run team and platform (recurring)

- Typically one engineer plus one operations owner per initiative
- Inference and token usage sized by use-case complexity against external benchmarks
- Platform and vendor contracts (for example telephony minimums) do not flex with volume

### What we expect to ask for

- Dedicated build capacity per big bet, funded rather than side-of-desk, starting with the well-trodden use cases
- An agreed run model (engineer plus operations owner) for each initiative before scale
- A technology capacity view (Srini) that shows what can start in Q4 2026 versus H1 2027

# We propose to sequence on five criteria and to tier investment to confidence: proven use cases move first, unproven ones earn scale funding through pilots

| Criterion                                   | What we assess                                                                      | Evidence we use at S2/S3                                                                             | How it drives sequencing                                                                       |
|---------------------------------------------|-------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------|
| <b>1 Confidence in the use case</b>         | Is this a proven pattern at scale elsewhere, or a hypothesis for us?                | External reference cases; tech POC view of integration effort; pilot results (S4)                    | Proven: build now. Unproven: fund design and pilot only, scale on evidence                     |
| <b>2 Timing of value</b>                    | When does run-rate value land against the committed profile, especially H1 2027?    | Value-realization profile by quarter agreed with workstreams; pressure-tested with Finance next week | Pull left initiatives that can land in H1 2027; compress build durations where tech POCs agree |
| <b>3 Investment intensity and treatment</b> | One-time build versus recurring run cost; capex versus opex; central pool versus BU | Cost sizing by phase and role; token cost by complexity; controllership rules                        | Tier investment to confidence; watch initiatives whose recurring cost erodes gross savings     |
| <b>4 Dependencies and double counting</b>   | Overlap with role ratios, contract-worker lever, Wave 1 actions and other big bets  | Parent/child structure in Wave; common baselines per gate                                            | One number per lever; sequence enablers (for example PDLC) after the actions they depend on    |
| <b>5 Execution capacity</b>                 | Can the team actually build and absorb this given post-Wave-1 capacity?             | Tech capacity view (Srini); spend to date versus spend requested                                     | Do not fund beyond demonstrated ability to spend; add kill criteria at each gate               |

## Tiered investment

Release funding in tranches tied to confidence and gate evidence, not to the full business case up front

## One holistic plan

Stack S2 cases across all levers into a single implementation plan at S3 with no double counting

## Triangulate, then choose

The investment pool is finite: weigh financial commitments, customer (NPS) goals and available dollars together


## Fail fast

Explicit continue, pivot or stop criteria at each gate, so a bet that will not deliver is stopped early and its funding redirected



# Applied to today's pipeline, the criteria point to three horizons: push proven use cases left now, pilot the rest through H1 2027, and gate scale on evidence

Illustrative placement by the ETO working team as of Sep 2, for discussion; to be confirmed with sponsors and at S3 once business cases are stacked across levers



# Investment will be material and front-loaded into build, with recurring run cost as the swing factor; sized numbers follow the Finance review next week

Observations across the portfolio; individual initiative figures are deliberately not shown until sponsors and Finance have reviewed them

## 1 Magnitude

The ask will be material and multi-year: one-time build across roughly 15-16 initiatives plus recurring run cost. We will bring ranges next week, not point estimates.

## 2 Shape of the cost

One-time cost is people by phase (plan and design, build and test, pilot, scale). Recurring cost is run teams plus inference; for SDLC recurring inference is the largest line and can erode gross savings.

## 3 Where the value comes from

Within servicing, value is dominated by contact avoidance rather than agent efficiency. How seller-productivity value is taken, cost-out versus reinvestment into revenue, is an open question with the sponsor and Finance.

## 4 Accounting treatment

Capitalization is decided at role level and by build type (from-scratch internal builds versus enhancements to vendor products). Rules and the amortization period to be confirmed with Controllership.

## 5 Funding source

Open question of how much is funded through the central pool versus BU budgets, and the capex/opex split. Build spend is assumed to start Q4 2026; an earlier start is being tested if treatment allows.

## 6 Scale and timing outliers

A small number of cases are on a different scale from the rest of the portfolio and should be sequenced to demonstrated ability to spend; PDLC value ramps late by design. Detail in the talk track, not on the page.

Implication: budget for tranches released against gate evidence; initiative-level figures follow once S2 submissions close and sponsors and Finance have reviewed them.

# Finance and HR are now inside the business cases, and accounting treatment is being worked with Controllership so that S3 cases arrive with a defensible funding view

## How Finance and HR are engaged

| Partner                      | Role in the business cases                                                                                                                            | Next touchpoint                                                                                  |
|------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------|
| Finance (Tom and team)       | Gives input as S2 cases land; challenges the value-realization profile by quarter, the ramp and the assumptions                                       | Finance input this week; review and challenge next week feeding the Jaime pressure-test sessions |
| HR (John, Alex)              | Translates milestones into headcount and partial capacity; to plan org-design implications (supervisory and management structures) as work is removed | Receives full models via ETO SharePoint; joins the Jaime sessions                                |
| Controllership               | Rules on capitalization by job family and build type; amortization period                                                                             | Sit-down to be scheduled; target a ruling before S3                                              |
| Technology points of contact | Roles per build, phase durations, what can be compressed or parallelized                                                                              | Build-duration challenge sessions before S3                                                      |

## Accounting treatment: working principles

- **Capitalization is decided at role level, not initiative level**
  - Engineering and similar build roles typically qualify; product management does not
- **Build type matters**
  - From-scratch internal builds and enhancements to a vendor product are treated differently
- **Amortization**
  - Models carry a working assumption on the period; Controllership to confirm
- **Benefits**
  - Not capitalized; recognized linearly as realized
- **Timing**
  - Build spend assumed from Q4 2026; an earlier start is being tested if treatment allows


## Open questions we are working before S3

- Share of investment funded from the central pool versus BU budgets, and the capex/opex split once role-level capitalization is ruled
- Whether an earlier build start accelerates enough benefit into Q1 2027 to justify the earlier spend, and how shared contract minimums (for example telephony) are allocated as contact volume falls

# Follow-ups from the CEO/CFO readout and the Finance model review centre on pace: the committed profile needs value in H1 2027, so we are shifting proven work left

Items being worked with Finance and HR after the readout; list and status to be confirmed against the readout notes with Gilad and Tawanda before the ELT session

Shape of the big-bet value build-up (illustrative, not to scale; profile under Finance review)



| Follow-up                       | What we are doing                                                                                      |             | ETO comment                |
|---------------------------------|--------------------------------------------------------------------------------------------------------|-------------|----------------------------|
| Ramp is too slow for H1 2027    | Faster scenario for servicing (agentic voice) and SDLC; challenge build durations with tech POCs       | <div></div> | In progress with tech POCs |
| Investment timing               | Test an earlier build start; confirm treatment with Controllershship; Finance reviewing hiring timing  | <div></div> | Depends on treatment       |
| Confidence and contingency      | Tier investment to confidence; keep the bottom-up pipeline filling in parallel; explicit gate criteria | <div></div> | For S3                     |
| Platform-scale investments      | Review spend plans against spend to date before sequencing                                             | <div></div> | Review with Finance        |
| Seller productivity framing     | Finalize the cost-out versus reinvestment framing with the sponsor and Finance                         | <div></div> | Sponsor review Sep 3       |
| Recurring inference cost (SDLC) | Token-efficiency approach with the SDLC lead so recurring cost does not erode gross savings            | <div></div> | In progress                |
| Fail-fast discipline            | Define continue, pivot or stop criteria per gate so a bet that will not deliver is stopped early       | <div></div> | For S3                     |

Bottom line: Finance can live with the ramp only if the total lands within the horizon and H1 2027 moves earlier; every option we are working either moves proven value earlier or moves spend later.

# Servicing sizing is aligned with ops leaders and seller productivity goes to its sponsor this week; PDLC and risk are in validation, and pilots begin at S4


Status is the ETO read as of Sep 2; the proof-of-concept column is to be confirmed with ETO leadership before the ELT session

| Area                |                                                                                    | What happened since the last update                                                                                                                                                                                            | What is next                                                                                                                                                                      | Proof point / pilot (to confirm)                                                                                                                             |
|---------------------|------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Servicing           |   | Model aligned with ops leaders on Sep 2 across contact avoidance, enhanced chat and self-service, agentic voice, skill-based routing and workforce management; copilot benefit folded into avoidance; non-FTE costs simplified | S2 submission this week; faster scenario for agentic voice; top-of-funnel headroom (help center, self-resolve) and support-function and leadership-span efficiencies sized for S3 | Pilot design at S4; NPS impact stated qualitatively (expected constant to slightly improved); data-driven method ready when the volume forecast is available |
| Seller productivity |   | Three fast-tracked initiatives (underwriting, pricing, SMB long tail) pressure-tested with the sales working team; S1-to-S2 deltas prepared for the sponsor review on Sep 3                                                    | Sponsor and Finance review; size remaining components (SMB long-tail attach, contra) after SME input                                                                              | Cost-out versus reinvestment framing to be finalized with sponsor and Finance                                                                                |
| SDLC                |   | Sizing with tech POCs; recurring inference cost identified as the key sensitivity                                                                                                                                              | Token-efficiency approach; build-duration challenge                                                                                                                               | Well-trodden use case; early pilot candidate for push-left                                                                                                   |
| PDLC                |   | Agreed with checkout to run as one centrally managed parent initiative; BU placeholder zeroed pending validation on re-cut baselines                                                                                           | Re-baseline with role-ratio actions; deeper dive with BU teams once numbers are ready                                                                                             | Design-and-pilot phase before any scale decision; value builds after the first year                                                                          |
| Risk                |  | Collections submitted for S2; financial-crimes cases prepared for Finance calls; Helix case handled on its own track                                                                                                           | Finance calls this week; Helix savings case still being clarified and investment request to be reviewed with Finance                                                              | Pilots for the smaller cases at S4; platform case sequenced after its execution plan is reviewed                                                             |

Read-across: sizing alignment with workstream leaders is the proof point available before S4; formal proofs of concept and pilots start once S3 fixes scope and funding.

# Next two weeks: close S2, run the Finance and Jaime reviews, and stack the cases into one S3 plan; we need your alignment on the criteria and a path on resourcing

## Timeline



## What we need from you



# Appendix

- A1 Stage-gate definitions
- A2 Cost-sizing methodology
- A3 Servicing big-bet architecture and NPS approach
- A4 Seller productivity: from S1 long list to three fast-tracked S2 cases
- A5 PDLC: one centrally managed parent initiative
- A6 Mapping of this document to the Friday and Tuesday agendas
- A7 Glossary

# Stage gates S1 to S5 move each initiative from a theoretical sizing to committed scale funding; automation initiatives are inherently iterative across these gates

Working definitions used by the ETO working team; to be confirmed against the ETO stage-gate process document

| Gate | Purpose                                                | Inputs                                                                                         | Outputs                                                           | Confidence in numbers    |
|------|--------------------------------------------------------|------------------------------------------------------------------------------------------------|-------------------------------------------------------------------|--------------------------|
| S1   | Identify and size the long-list of opportunities       | Outside-in benchmarks; initial baselines; sponsor interviews                                   | Initiative list with preliminary value and investment             | Directional; ranges wide |
| S2   | Refine business cases with SMEs, data and technology   | SME interviews; actual data at intent or role level; tech POC roles and phases; Finance review | Submitted business cases in Wave with one-time and recurring cost | Higher; still pre-pilot  |
| S3   | Stack cases across levers into one implementation plan | All S2 cases; parent/child structure; sequencing criteria; capacity view                       | Holistic plan; sequencing; resourcing and funding request         | Best theoretical view    |
| S4   | Prove through pilots and proofs of concept             | Pilot results; adoption; real build durations; org-design ripple effects                       | Evidence-based timelines; kill or scale decisions                 | Evidence-based           |
| S5   | Commit scale funding and targets                       | Updated sizing on pilot evidence; treatment confirmed                                          | Committed run-rate and investment by initiative                   | Committed                |

### Why numbers move between gates

Baselines, SME input and technology effort are refined at each gate; S4 pilots replace assumptions with evidence

### What is fixed at S3

Scope, sequencing, resourcing plan and the funding request, with continue, pivot or stop criteria per initiative

### What is fixed at S5

Committed run-rate and investment by initiative once pilot evidence and accounting treatment are in (to confirm against the ETO process)



# Investment is sized bottom-up as people cost by phase plus non-people run cost, using a typical build team and token benchmarks by use-case complexity

## One-time cost: people by phase

- Standard build team: product management, engineering, QA, DevOps, UX
- Blended cost per person applied across the team
- Work broken into phases: plan and design, build and test, pilot, scale
- For each phase, number of teams or people and duration agreed with technology POCs
- Parallelization assumed where phases can overlap; build-duration compression being tested for the faster scenario
- Capitalization assessed at role level and by build type; rules to be confirmed with Controllership

## Recurring cost: run team and inference

- Run team per initiative, typically one engineer plus one operations owner
- Token and inference usage sized by defining the complexity of each use case
- Industry benchmarks for average token usage at that complexity applied per use case
- SDLC carries the highest inference load (tokens consumed continuously while writing code); efficiency work under way
- Platform and vendor contracts with monthly minimums (for example telephony) do not fall with volume; how shared contract cost is allocated across initiatives is still to be agreed
- Benefits are not capitalized and are recognized linearly as realized

## Working assumptions currently in the models (to be confirmed at the Finance review)

### Build start

All initiatives assumed to start build in Q4 2026; an earlier start is being tested where treatment allows

### Amortization

Working assumption on the period for capitalizable roles; Controllership to confirm

### Benefit haircuts

Conservative haircuts applied to benefit assumptions, with upside noted rather than sized

### Run cost

Token usage per use case from complexity benchmarks; SDLC the largest recurring line

*Outputs feed Wave as one-time and recurring cost per initiative; net run-rate benefit in Wave must include both, with assumptions noted for the Finance review.*

# The servicing big bet works down the contact funnel: avoid the contact, self-serve it, resolve it with an AI agent, then run the remaining human queue more efficiently

## 1 Avoid the contact

Policy, product and journey fixes so the need never arises; top-of-funnel headroom (help center, surfacing known data) flagged for S3

## 2 Self-serve it

Enhanced chat and self-service on the post-avoidance baseline, with containment capped by complexity

## 3 Resolve with an AI agent

Agentic voice takes contacts end to end; assistant-style benefits folded in rather than counted separately

## 4 Run the rest efficiently

Skill-based routing and workforce management; non-FTE costs (telephony, software) simplified


### NPS position (ETO working view, not yet reviewed with ops leaders)

Expected constant to slightly improved across the levers; a quantified waterfall is a separate exercise for S3 and needs the 18-month contact forecast and headcount plan

### Guardrails agreed with ops leaders

Conservative sizing in Wave with upside noted; one line per lever; no separate efficiency claim that depends on unproven tooling; initiatives that generate contacts carry that cost in the momentum case

# Seller productivity moved from the mid-August long list to three fast-tracked S2 cases refined with the sales working team over three weeks; two inputs remain open



## What the three initiatives do

- Underwriting: automate the analysis and decision support behind merchant underwriting
- Pricing: automate pricing analysis and proposal preparation for sellers
- SMB long tail: agent-assisted outreach so partners and sellers can serve the long tail of merchants
- All three release seller and support capacity; how that capacity is taken (cost-out or reinvestment) is to be agreed with the sponsor and Finance


## Open items carried into S2

- SMB long-tail attach component: sizing pending SME interviews and data
- Contra-revenue effect: assumption not yet validated with the contra team
- Cost-out versus reinvestment framing to be agreed with the sponsor before Finance review
- A separate sales integration cost-out runs as a bottom-up initiative, not as a big bet

Detail on revenue, cost and one-time investment stays in the sponsor and Finance materials until S2 is submitted and reviewed.

# PDLC is now one centrally managed big bet with a company-level parent initiative and BU child views, re-baselined for Wave 1 and role-ratio actions and phased design-first

## Structure in Wave



## Phased, design-first roadmap

### 1 Define and prove

Define the target lifecycle, technology backbone and KPIs; run one or two pilots to prove value in detail

### 2 Build across teams

Extend to a handful of teams across a couple of BUs; harden the agentic tooling, governance and change model

### 3 Scale and improve

Roll out broadly; value ramps only after stabilization, so little is expected in the first year

# One storyline serves both sessions: Friday uses the full document; Tuesday's ELT update reuses the pages mapped to its four agenda items with the pipeline snapshot refreshed

| Agenda item                              | Friday (Enrique, Jaime, Anshu)                                   | Tuesday ELT                                   | Pages            |
|------------------------------------------|------------------------------------------------------------------|-----------------------------------------------|------------------|
| Snapshot of initiatives and where we are | Level-set; status with ETO commentary                            | Overall pipeline tracking                     | 4, 5             |
| Org change impacts                       | Talk track                                                       | Context only                                  | 6                |
| How we are resourcing; tech capacity     | Core topic; Srini framework page to insert                       | Approach to resourcing                        | 7 (+ Srini page) |
| Sequencing and prioritization            | Proposed criteria and illustrative horizons                      | Approach to sequencing and prioritization     | 8, 9             |
| Investment themes; accounting treatment  | Observations and themes; sized numbers follow the Finance review | Not covered until Finance review              | 10, 11           |
| CEO/CFO readout follow-ups               | Preview                                                          | Follow-ups from CEO/CFO readout               | 12               |
| Progress and proofs of concept           | Reference                                                        | Progress and status on big bets including POC | 13               |
| Next steps and asks                      | Asks to Enrique and Jaime                                        | Next steps                                    | 14               |

# Glossary of terms used in this document

| Term              | Meaning                                                                                             | Term                      | Meaning                                                                                                                                |
|-------------------|-----------------------------------------------------------------------------------------------------|---------------------------|----------------------------------------------------------------------------------------------------------------------------------------|
| Big bet           | Large, centrally tracked automation initiative; roughly 15-16 across five areas                     | Wave                      | Initiative tracking tool for business cases; also Wave 1 (org actions completed) versus Wave 2                                         |
| S1 to S5          | Stage gates from initial sizing to committed scale (definitions to confirm against the ETO process) | Tech POC                  | Technology point of contact per initiative (distinct from proof of concept)                                                            |
| One-time cost     | Build cost: people by phase plus set-up                                                             | Recurring cost            | Run team (typically one engineer plus one operations owner) plus inference and token usage; shared vendor contracts tracked separately |
| Contact avoidance | Removing the need for a customer contact entirely                                                   | Containment               | Share of contacts resolved in a digital channel without a human                                                                        |
| Agentic voice     | AI agent that handles voice contacts end to end                                                     | Copilot                   | AI assistant that supports a human agent during a contact                                                                              |
| PDLC / SDLC       | Product / software development lifecycle                                                            | Parent / child initiative | Structure in Wave that counts a lever once at company level with BU views beneath                                                      |
| Central pool      | Company-level funding for transformation investment, versus BU budgets                              | Run-rate value            | Annualized gross benefit once an initiative is fully ramped                                                                            |